

Assignment3

Gourp member: Mike Wang, Ying Zhou, Hanjie(Huntley) Liu

Kaggle marks account for part3: Mike Wang

Hanjie Liu

Student ID: 301404949

Kaggle Name: Huntley Liu

Late day: 1

Part1

List of configs

```
...
# Set the configs for the detection part in here.
# TODO: approx 15 lines
...
cfg = get_cfg()
cfg.OUTPUT_DIR = "{}/output/".format(BASE_DIR)

cfg.merge_from_file(model_zoo.get_config_file("COCO-InstanceSegmentation/mask_rcnn_R_50_FPN_3x.yaml"))
#cfg.merge_from_file(model_zoo.get_config_file("COCO-Detection/faster_rcnn_X_101_32x8d_FPN_3x.yaml"))
cfg.DATASETS.TRAIN = ("plane_train",)
cfg.DATASETS.TEST = ("plane_test",)

cfg.DATALOADER.NUM_WORKERS = 4
cfg.MODEL.WEIGHTS = model_zoo.get_checkpoint_url("COCO-InstanceSegmentation/mask_rcnn_R_50_FPN_3x.yaml")
#cfg.MODEL.WEIGHTS = model_zoo.get_checkpoint_url("COCO-Detection/faster_rcnn_X_101_32x8d_FPN_3x.yaml")
cfg.SOLVER.IMS_PER_BATCH = 2
cfg.SOLVER.BASE_LR = 0.00025
cfg.SOLVER.MAX_ITER = 500
cfg.MODEL.ROI_HEADS.BATCH_SIZE_PER_IMAGE = 128
cfg.MODEL.ROI_HEADS.NUM_CLASSES = 1
```

List of modifications and Improving factors:

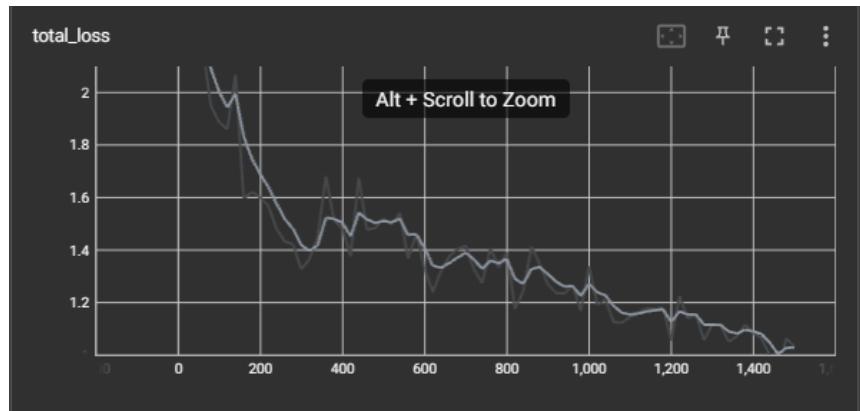
1. The Base_Lr has been modified from 0.001 to 0.00025. When I using 0.001 as my Base_Lr, the loss goes down very fast at beginning but it is hard to converge in last 100 iterations. So I change it to 0.00025, the accuracy steadily improved. I found that the higher Base_Lr will have a lower accuracy and lower Base_Lr will take longer time to converge.
2. I used "faster_rcnn_X_101_32x8d_FPN_3x.yaml" from Model_ZOO website as my pre-training model, it causes more time than tutorial's but has higher accuracy.
3. I change the iterations from 500 to 1500(5000 for part3), because 500 iterations still has fast speed of decreasing loss, but 1500 iterations become stable, and results are credible.

Training loss on the last part of iterations:

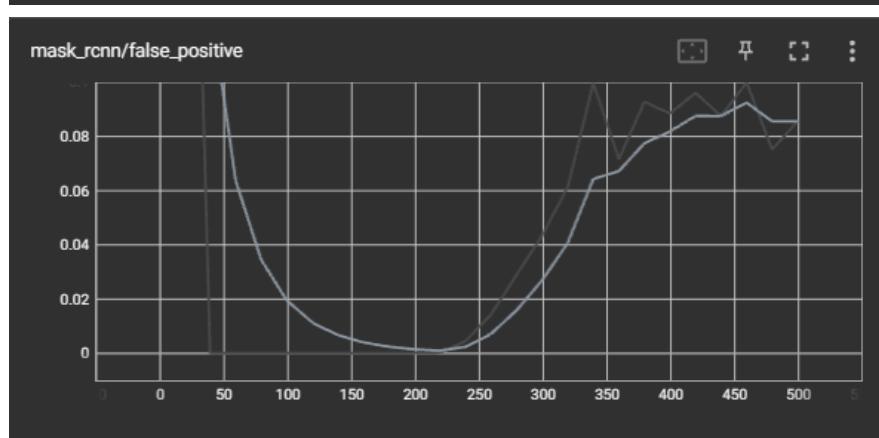
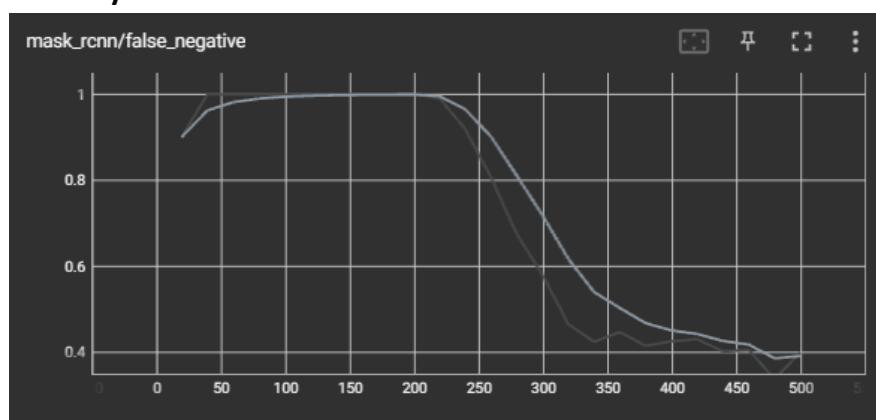
```
eta: 0:07:05 iter: 1319 total_loss: 1.114 loss_cls: 0.2389 loss_box_reg: 0.5352 loss_rpn_cls: 0.1155
eta: 0:06:18 iter: 1339 total_loss: 1.051 loss_cls: 0.2438 loss_box_reg: 0.5174 loss_rpn_cls: 0.09838
eta: 0:05:30 iter: 1359 total_loss: 1.07 loss_cls: 0.2429 loss_box_reg: 0.5092 loss_rpn_cls: 0.09847
eta: 0:04:43 iter: 1379 total_loss: 1.115 loss_cls: 0.2425 loss_box_reg: 0.5291 loss_rpn_cls: 0.08864
eta: 0:03:56 iter: 1399 total_loss: 1.081 loss_cls: 0.2768 loss_box_reg: 0.4954 loss_rpn_cls: 0.09134
eta: 0:03:09 iter: 1419 total_loss: 1.064 loss_cls: 0.2567 loss_box_reg: 0.4988 loss_rpn_cls: 0.08788
eta: 0:02:22 iter: 1439 total_loss: 0.9969 loss_cls: 0.2459 loss_box_reg: 0.4948 loss_rpn_cls: 0.0919
eta: 0:01:34 iter: 1459 total_loss: 0.9374 loss_cls: 0.2385 loss_box_reg: 0.5003 loss_rpn_cls: 0.0853
eta: 0:00:47 iter: 1479 total_loss: 1.062 loss_cls: 0.2201 loss_box_reg: 0.5333 loss_rpn_cls: 0.1102
eta: 0:00:00 iter: 1499 total_loss: 1.032 loss_cls: 0.2396 loss_box_reg: 0.4709 loss_rpn_cls: 0.1092
```

Visualization of test samples

Total loss:



Accuracy:



Test sample 1 for iter 500:



Test sample 2 for 500 iter:



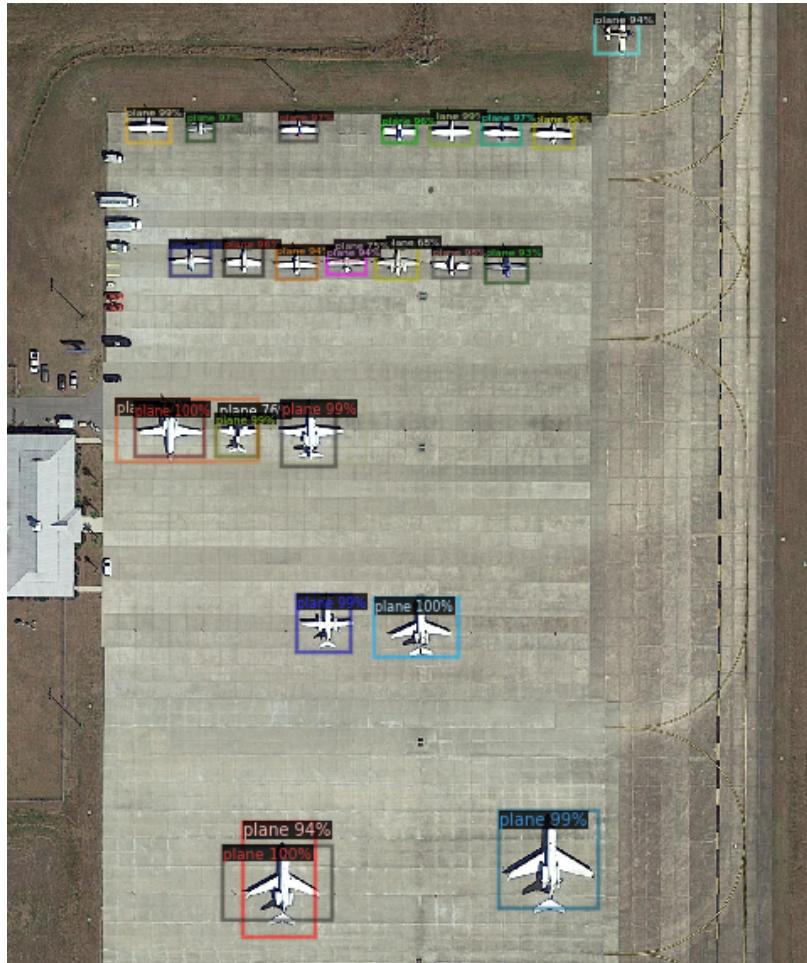
Test sample 3 for 500 iter:



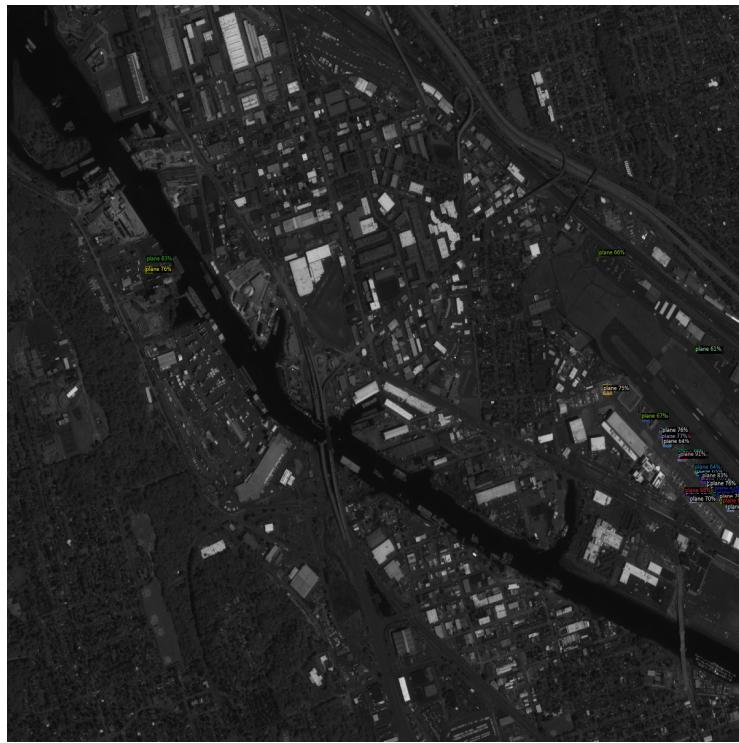
Test sample 1 for 1500 iter:



Test sample 2 for 1500 iter:



Test sample 3 for 1500 iter:



Final accuracy for the AP50 on 500 iter:

```
[02/28 21:00:44 d2.evaluation.fast_eval_api]: Evaluate annotation type *bbox*
[02/28 21:00:44 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finished in 0.26 seconds.
[02/28 21:00:44 d2.evaluation.fast_eval_api]: Accumulating evaluation results...
[02/28 21:00:44 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finished in 0.02 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.317
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.532
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.351
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.220
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.404
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.618
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.015
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.130
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.357
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.222
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.447
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.738
[02/28 21:00:44 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APM | API |
|-----:|-----:|-----:|-----:|-----:|-----:|
| 31.708 | 53.240 | 35.065 | 22.005 | 40.373 | 61.828 |
Loading and preparing results...
DONE (t=0.15s)
creating index...
index created!
[02/28 21:00:45 d2.evaluation.fast_eval_api]: Evaluate annotation type *segm*
[02/28 21:00:46 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finished in 0.94 seconds.
[02/28 21:00:46 d2.evaluation.fast_eval_api]: Accumulating evaluation results...
[02/28 21:00:46 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finished in 0.02 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.098
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.328
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.022
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.053
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.114
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.336
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.008
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.060
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.129
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.070
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.157
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.396
[02/28 21:00:46 d2.evaluation.coco_evaluation]: Evaluation results for segm:
| AP | AP50 | AP75 | APs | APM | API |
|-----:|-----:|-----:|-----:|-----:|-----:|
| 9.835 | 32.806 | 2.193 | 5.293 | 11.432 | 33.554 |
OrderedDict([('bbox', {'AP': 31.70845464725845, 'AP50': 53.239532202057106, 'AP75': 35.06519663659673,
```

Final accuracy for the AP50 on 1500 iter:

```
[03/02 06:26:01 d2.evaluation.fast_eval_api]: Evaluate annotation type *bbox*
[03/02 06:26:02 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finished in 0.29 seconds.
[03/02 06:26:02 d2.evaluation.fast_eval_api]: Accumulating evaluation results...
[03/02 06:26:02 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finished in 0.03 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.287
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.585
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.247
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.231
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.349
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.409
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.014
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.115
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.338
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.244
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.407
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.558
[03/02 06:26:02 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APM | API |
|-----:|-----:|-----:|-----:|-----:|-----:|
| 28.694 | 58.545 | 24.709 | 23.135 | 34.925 | 40.933 |
OrderedDict([('bbox', {'AP': 28.694165524522607, 'AP50': 58.545235702691336, 'AP75': 24.708872603354685,
```

Ablation study:

Effect of increasing maximum iterations from 500 to 1500: We increase our iteration round from 500 to 1500, in this situation, the accuracy of object are impressively increase. Some of them has been improved from 70% to 99%, and the program can detect more planed as it predictor probability exceed 0.6. The AP50 detection accurate also gose up from 53 to 58. So we know that more iteration before overfit will increase our accuracy

Part 2

Hyperparameter setting:

```
# Set the hyperparameters
num_epochs = 50
batch_size = 4
learning_rate = 0.001
weight_decay = 1e-5
```

The optimizer: "torch.optim.SGD"

Net architecture:

I add four more down and up sampling layers in network. The input layer was given (3, 16) and each down sampling layer increase the output to the twice of input, until we got 512. The up sampling will decrease the output to half of the input until we got 4. We increase the feature channels for helping us detecting feature more credible. The loss is going stable around 0.068.

```
def __init__(self):
    super(MyModel, self).__init__()

    self.input_conv = conv(3, 16)
    self.down1 = down(16, 32)
    self.down2 = down(32, 64)
    self.down3 = down(64, 128)
    self.down4 = down(128, 256)
    self.down5 = down(256, 512)

    self.up1 = up(512, 256)
    self.up2 = up(256, 128)
    self.up3 = up(128, 64)
    self.up4 = up(64, 32)
    self.up5 = up(32, 4)
    self.output_conv = conv(4, 1, False) #

def forward(self, input):
    y = self.input_conv(input)
    y = self.down1(y)
    y = self.down2(y)
    y = self.down3(y)
    y = self.down4(y)
    y = self.down5(y)
    y = self.up1(y)
    y = self.up2(y)
    y = self.up3(y)
    y = self.up4(y)
    y = self.up5(y)
    output = self.output_conv(y)
    return output
```

Loss function:

```
nn.BCEWithLogitsLoss()
```

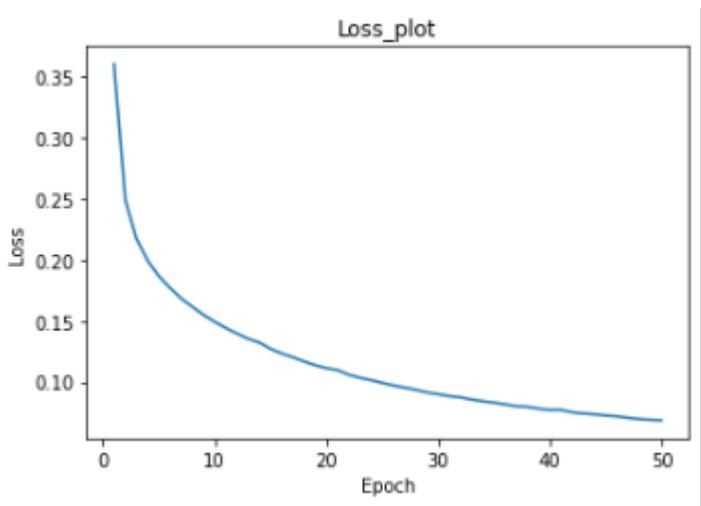
Traning loss of first 3 epochs:

```
<ipython-input-35-5807c010b2c5>:25: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, device=torch.device('cuda'), requires_grad = True)
<ipython-input-39-5807c010b2c5>:25: UserWarning: To copy construct from a tensor (t
  mask = torch.tensor(mask, device=torch.device('cuda'), requires_grad = True)
Epoch: 0, Loss: 0.35984018445014954
100% [2000/2000 [00:58<00:00, 37.52it/s]
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
Epoch: 1, Loss: 0.2489643692970276
100% [2000/2000 [00:59<00:00, 37.74it/s]
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
Epoch: 2, Loss: 0.21722055971622467
100% [2000/2000 [00:57<00:00, 35.11it/s]
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
```

Traning loss of last 3 epochs:

```
100% [2000/2000 [00:59<00:00, 36.65it/s]
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
Epoch: 47, Loss: 0.06972382217645645
100% [2000/2000 [00:58<00:00, 37.43it/s]
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
Epoch: 48, Loss: 0.06918254494667053
100% [2000/2000 [00:58<00:00, 36.46it/s]
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
<ipython-input-35-8ff7c3079b41>:29: UserWarning: To copy construct from a tensor (t
  img = torch.tensor(img, dtype=torch.float)
Epoch: 49, Loss: 0.06867346167564392
```

Plot of loss:



Final mean IoU: **IoU: 0.9006885895082279**

Test sample: Original; predicted; test

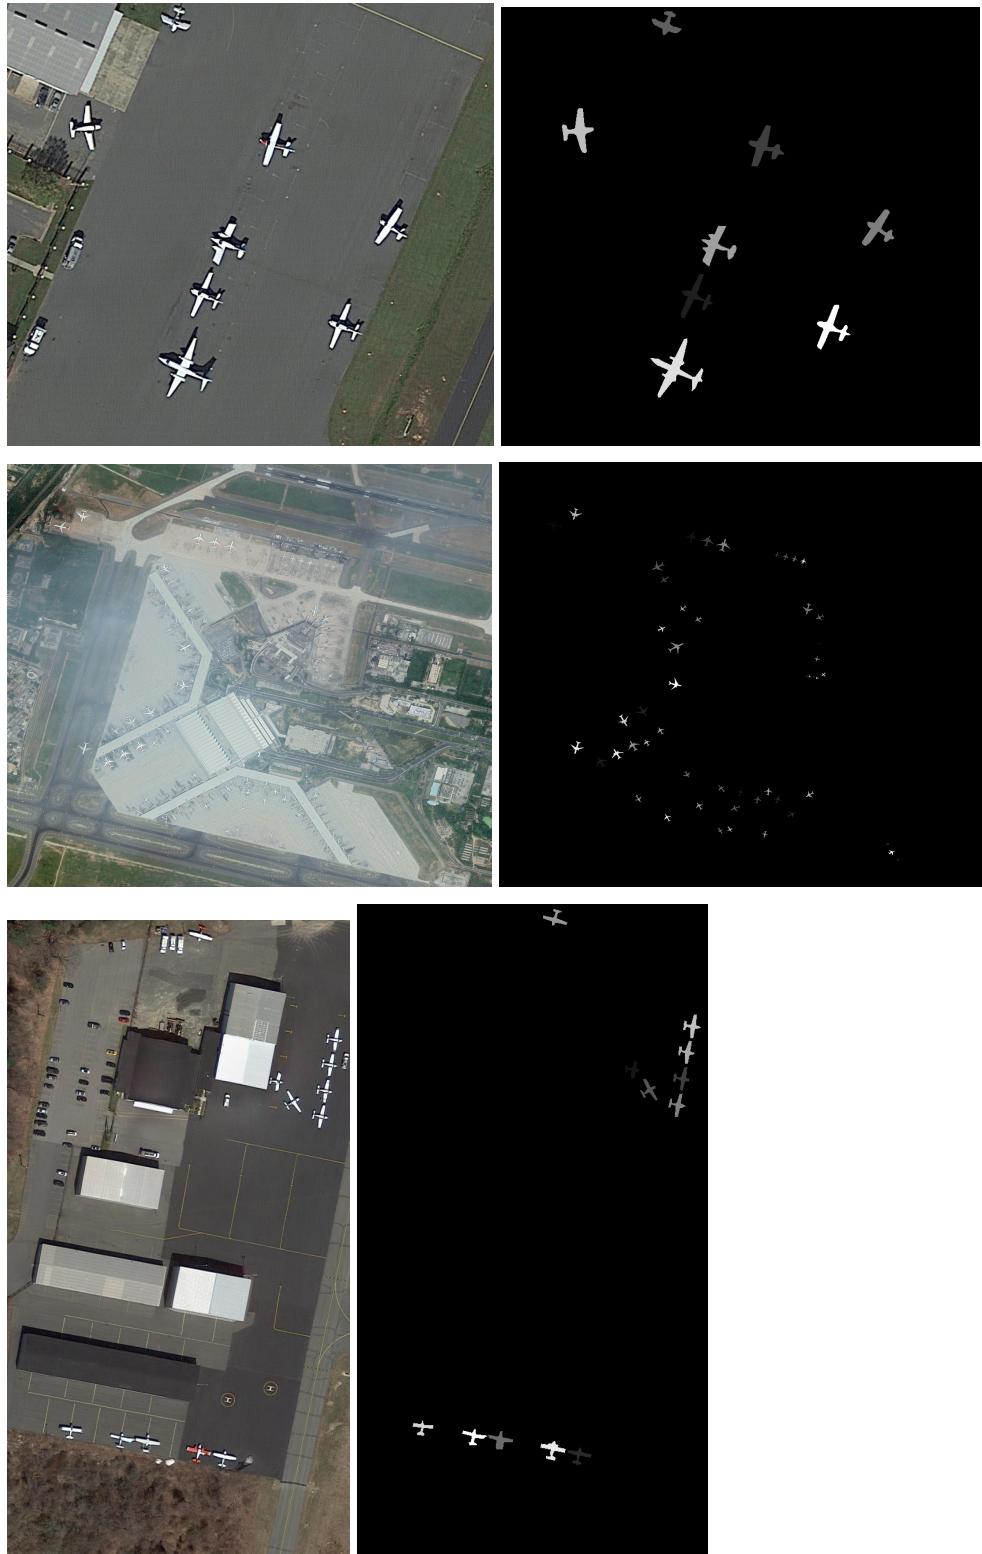


Part 3:

Kaggle ID: Mike Wang (uploader)

Best score: **0.46629**

Test and Train sample:

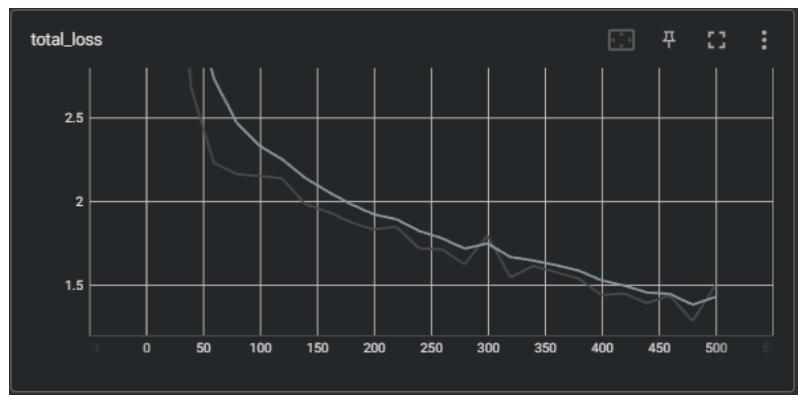


Part 4:

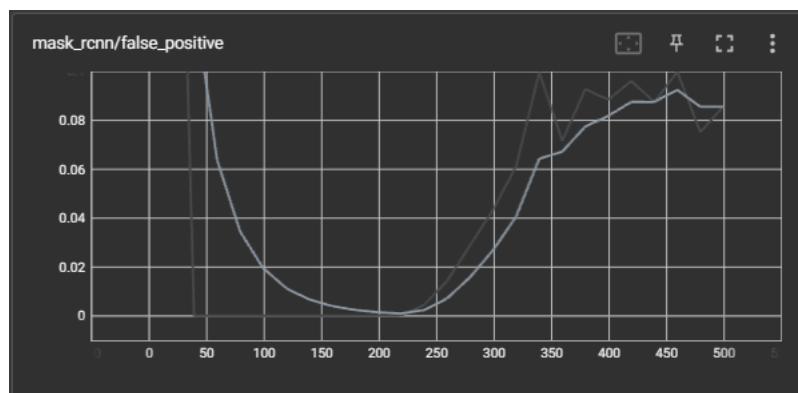
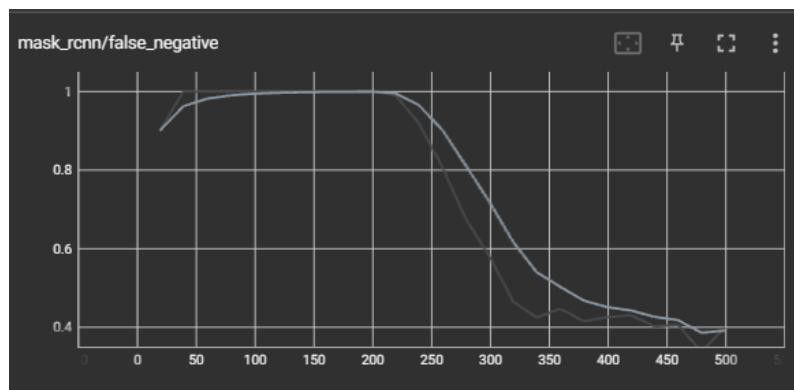
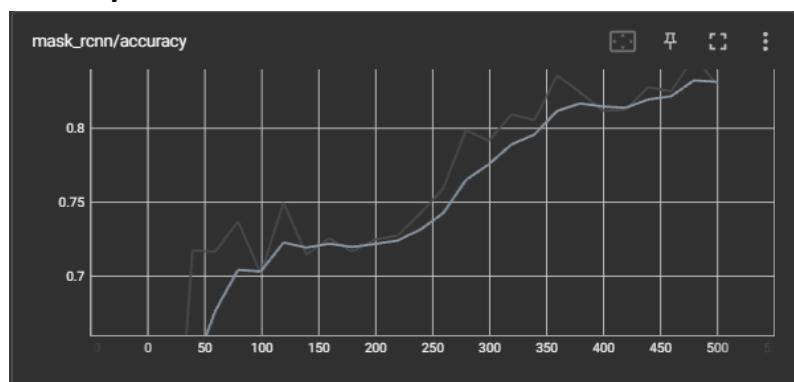
Our model is “COCO-InstanceSegmentation/mask_rcnn_R_50_FPN_3x.yaml”

We run 100 iterations and 500 iterations:

Total loss:



Accuracy:



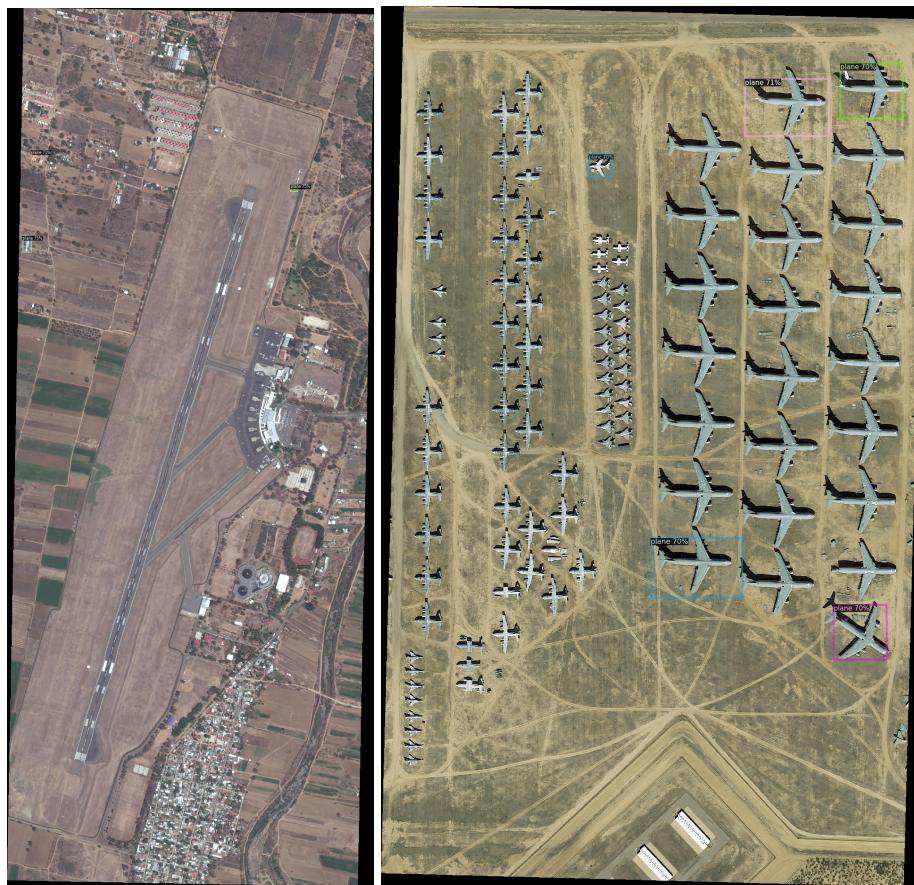
Accuracy of AP50:

```

index created:
[02/23 07:25:09 d2.evaluation.fast_eval_api]: Evaluate annotation type *bbox*
[02/23 07:25:09 d2.evaluation.fast_eval_api]: COCOeval.opt.evaluate() finished in 0.28 seconds.
[02/23 07:25:09 d2.evaluation.fast_eval_api]: Accumulating evaluation results...
[02/23 07:25:09 d2.evaluation.fast_eval_api]: COCOeval.opt.accumulate() finished in 0.03 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.249
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.467
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.246
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.166
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.330
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.506
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.013
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.113
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.296
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.173
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.378
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.650
[02/23 07:25:09 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | API |
|-----|-----|-----|-----|-----|-----|
| 24.939 | 46.740 | 24.626 | 16.622 | 32.966 | 50.553 |
Loading and preparing results...
DONE (t=0.16s)
creating index...
index created!
[02/23 07:25:10 d2.evaluation.fast_eval_api]: Evaluate annotation type *segm*
[02/23 07:25:11 d2.evaluation.fast_eval_api]: COCOeval.opt.evaluate() finished in 0.97 seconds.
[02/23 07:25:11 d2.evaluation.fast_eval_api]: Accumulating evaluation results...
[02/23 07:25:11 d2.evaluation.fast_eval_api]: COCOeval.opt.accumulate() finished in 0.03 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.054
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.214
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.007
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.030
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.058
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.218
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.006
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.041
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.080
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.041
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.094
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.274
[02/23 07:25:11 d2.evaluation.coco_evaluation]: Evaluation results for segm:
| AP | AP50 | AP75 | APs | APm | API |
|-----|-----|-----|-----|-----|-----|
| 5.406 | 21.382 | 0.716 | 2.980 | 5.849 | 21.831 |
OrderedDict([('bbox', {'AP': 24.9392739530302, 'AP50': 46.73957178167791, 'AP75': 24.62580693029191, 'APs': 16.622, 'APm': 32.966}), ('segm', {'AP': 5.406, 'AP50': 21.382, 'AP75': 0.716, 'APs': 2.98, 'APm': 5.849})])

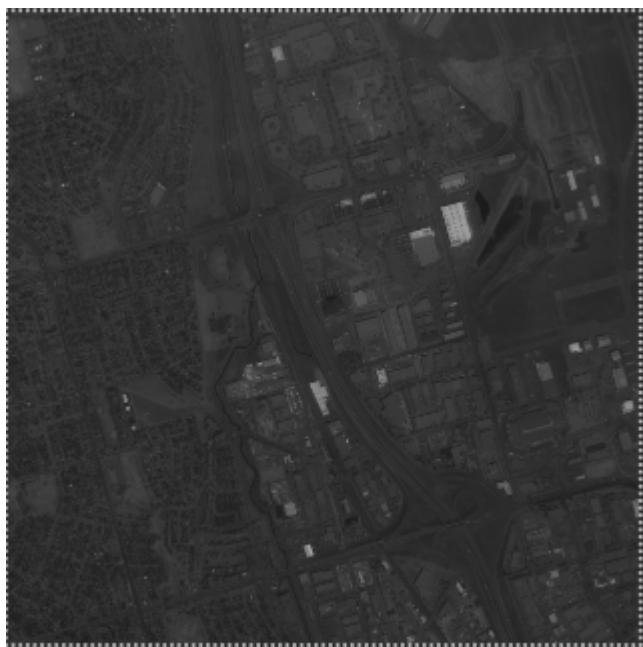
```

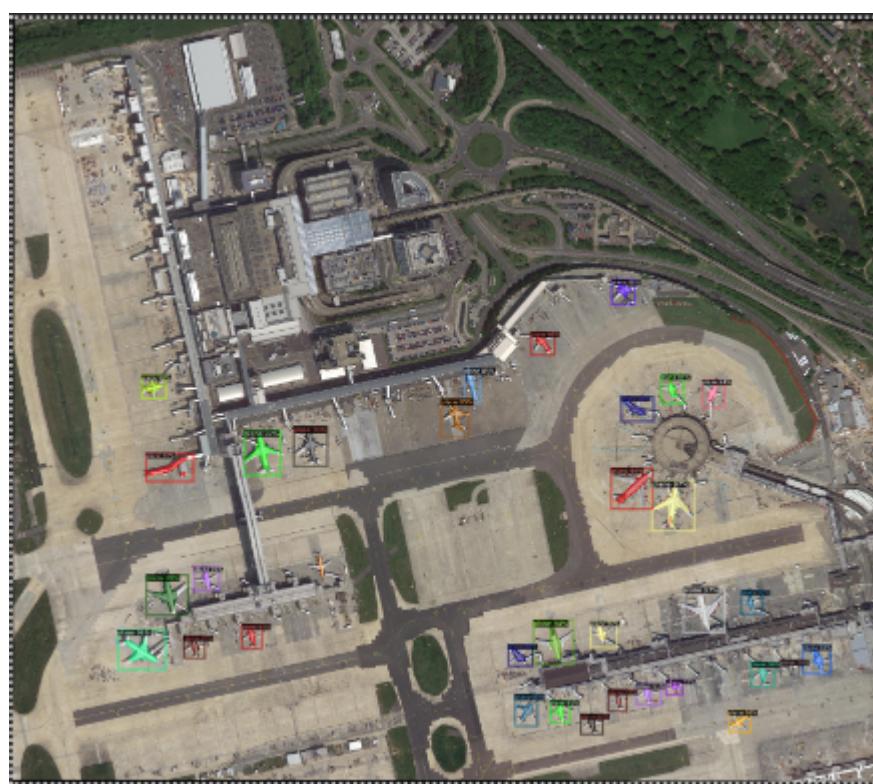
Test for 100 iter:





Test for iter 500:





Mean IOU:

```
IoU: 0.8160094420176907
```

Difference:

Compared both results, this model can be run less than previous one for same bbox accuracy, such as 500. The part 1 model does not work well in 500 iterations, but this one looks fine. And Its accuracy are increasing impressivly when iteration goes from 100 to 500. So we can say this model can be used faster than previous one.

From the evaluation output, the AP50 of this model for the bounding box is less than previous model, eventhough while they have same iterations. So due to AP50 of this one is lower than previous one, we can say this model is less accurate than previous one.